

Same runtime, different altitudes

The three things are not a product menu. They are one stack, and you can move between them without a rewrite.

BUILT HERE

LangChain

The framework. Building blocks and `create_agent` — a configurable model-and-tool harness with middleware.

WHEN YOU REACH FOR IT

Default starting point. Nearly every agent, including this one.

LangGraph

The runtime underneath. Durable execution, state, checkpointing, streaming, human-in-the-loop.

WHEN YOU REACH FOR IT

Directly, when you need explicit state transitions, deterministic routing, or parallel branches.

Deep Agents

An opinionated harness on top. Planning, subagents, virtual filesystem, skills.

WHEN YOU REACH FOR IT

Long-horizon work where the context outgrows the window.

You don't pick one. `create_agent` returns a compiled LangGraph. What I built today drops into a `StateGraph` as a node tomorrow with its middleware intact — no rewrite at the boundary.

LangSmith is the other half. The OSS is how you build the agent; LangSmith is how you find out whether it's any good. Framework-agnostic — it speaks OpenTelemetry.

RESULTS

Every number labeled

Measured on this build · Modeled from assumptions · Unavailable until real traffic. The third category is the one worth reading.

METRIC	VALUE	BASIS
Task resolution, benchmark 30 examples · 6 slices · every slice clean at this commit	30 / 30	MEASURED
Cross-tenant data access code evaluator, never a judge · hard gate	0 (30/30)	MEASURED
Unapproved writes no refund ticket without a human click	0 (30/30)	MEASURED
Escalation decisions correct unsafe-resolution rate 0 · 4 labeled cases, so a percentage would be dishonest	4 / 4	MEASURED
Full-turn latency, p50 p95 7.1s — directional only at n=30	3.8 s	MEASURED
Cost per priced eval turn 29 of 30 examples are single-turn · a real conversation costs several times more	\$0.0010	MEASURED
API spend, every eval run on disk 220 priced rows across 19 result files · both arms, every rerun · excludes Studio	\$0.22	MEASURED
Deterministic tests passing the security claims each have one behind them	139	MEASURED
Deflected labor / month 1,200 contacts · 60% billing-class · 70% containment · 6 min at \$0.60/min	~\$1,800	MODELED
Containment rate · recommendation attach rate the North Star and the upsell line — both need production traffic	not known	UNAVAILABLE

THE EXPERIMENT

Both arms resolved 30/30. I shipped flat.

Bars were pre-registered before the supervisor existed. It came in +2.8s on p50 against a +2.0s ceiling, with no mixed-intent gain and no quality gain to pay for it.

Last week this said 30 against 29, in the supervisor's favour, and I shipped flat anyway — one example on thirty isn't a quality difference. Rerun at this commit, they tie. That is what a one-example margin was worth. What reproduced both times is the overhead of the extra hop: 1.7x median latency, 61 tool calls against 45. The verdict rests on that.

PROVENANCE

Commit, both arms	a29a530 · clean
Agent model	gpt-5.6-luna
Dataset digest	6ba144216f8d
Evaluator digest	decda2664959

Generated into reports/experiment.md, not typed. The script refuses to publish two arms that disagree on commit, model, dataset, or evaluators.