

Same runtime, different altitudes

The three things are not a product menu. They are one stack, and you can move between them without a rewrite.

BUILT HERE

LangChain

The framework. Building blocks and `create_agent` — a configurable model-and-tool harness with middleware.

WHEN YOU REACH FOR IT

Default starting point. Nearly every agent, including this one.

LangGraph

The runtime underneath. Durable execution, state, checkpointing, streaming, human-in-the-loop.

WHEN YOU REACH FOR IT

Directly, when you need explicit state transitions, deterministic routing, or parallel branches.

Deep Agents

An opinionated harness on top. Planning, subagents, virtual filesystem, skills.

WHEN YOU REACH FOR IT

Long-horizon work where the context outgrows the window.

You don't pick one. `create_agent` returns a compiled LangGraph. What I built today drops into a `StateGraph` as a node tomorrow with its middleware intact — no rewrite at the boundary.

LangSmith is the other half. The OSS is how you build the agent; LangSmith is how you find out whether it's any good. Framework-agnostic — it speaks OpenTelemetry.

RESULTS

Every number labeled

Measured on this build · Modeled from assumptions · Unavailable until real traffic. The third category is the one worth reading.

METRIC	VALUE	BASIS
Task resolution, benchmark 30 examples · 6 slices · the one miss omitted the rep's name — content, not safety	29 / 30	MEASURED
Cross-tenant data access code evaluator, never a judge · hard gate	0 (30/30)	MEASURED
Unapproved writes no refund ticket without a human click	0 (30/30)	MEASURED
Escalation decisions correct unsafe-resolution rate 0 · 4 labeled cases, so a percentage would be dishonest	4 / 4	MEASURED
Full-turn latency, p50 p95 10.0s — directional only at n=30	4.0 s	MEASURED
Cost per priced eval turn 29 of 30 examples are single-turn · a real conversation costs several times more	\$0.0010	MEASURED
Total API spend, whole build OpenAI dashboard, not the estimate	\$ __. __	MEASURED
Deterministic tests passing the security claims each have one behind them	139	MEASURED
Deflected labor / month 1,200 contacts · 60% billing-class · 70% containment · 6 min at \$0.60/min	~\$1,800	MODELED
Containment rate · recommendation attach rate the North Star and the upsell line — both need production traffic	not known	UNAVAILABLE

THE EXPERIMENT

The supervisor won 30/30 against flat's 29. I shipped flat.

Bars were pre-registered before the supervisor existed. It came in +2.2s on p50 against a +2.0s ceiling, with no mixed-intent gain. One example out of thirty doesn't buy that.

One example does not establish a quality difference. Older runs at unrecorded commits put each arm at both 29 and 30, failing a different example each time — not controlled variance, but enough to show the count is sensitive. What is clear in the controlled comparison is the overhead of the extra hop: ~1.5x median latency, 70 tool calls against 46. The verdict rests on that, not on the one-example gap.

PROVENANCE

Commit, both arms	12cefb1 · clean
Agent model	gpt-5.6-luna
Dataset digest	6ba144216f8d
Evaluator digest	9e2516b82112

Generated into reports/experiment.md, not typed. The script refuses to publish two arms that disagree on commit, model, dataset, or evaluators.

Inference is not the binding cost. Even at the modeled full-conversation rate, deflecting 500 contacts is single-digit dollars of tokens against ~\$1,800 of labor. I'm not leading with that ratio — the real cost here is engineering and the review loop, not the model. 2 / 2